Tirth Snehal Patel – (NUID 002251943)

Big Data System Engineering with Scala Fall 2023 Spark Assignment No. 2



Below is the code implementation along with the output.

Section Spark Assignment 2

```
(https://databricks.com)
    %scala
    import org.apache.spark.sql.types.{StructType, StructField, StringType, IntegerType, FloatType};
    import org.apache.spark.sql.functions._
    import org.apache.spark.mllib.linalg.Vectors
    import org.apache.spark.ml.feature.VectorAssembler
    import org.apache.spark.ml.Pipeline
    import\ org. apache. spark. ml. classification. \{Random Forest Classification Model,\ Random Forest Classifier\}
    import\ org. apache. spark.ml. feature. \{IndexToString,\ StringIndexer,\ VectorIndexer\}
    import\ org. apache. spark. ml. evaluation. Binary Classification Evaluator
    import\ org. apache. spark. mllib. stat. \{ Multivariate Statistical Summary,\ Statistics \}
   import\ org. apache. spark. ml. evaluation. Multiclass Classification Evaluator
    import org.apache.spark.ml.tuning.{CrossValidator, ParamGridBuilder}
    import\ org. apache. spark. ml. classification. \{GBTClassificationModel,\ GBTClassifier\}
   import org.apache.spark.ml.classification.DecisionTreeClassificationModel
    import\ org. apache. spark.ml. classification. Decision Tree Classifier
  import org.apache.spark.sql.types.{StructType, StructField, StringType, IntegerType, FloatType}
  import org.apache.spark.sql.functions
  import org.apache.spark.mllib.linalg.Vectors
  import org.apache.spark.ml.feature.VectorAssembler
  import org.apache.spark.ml.Pipeline
  import org.apache.spark.ml.classification.{RandomForestClassificationModel, RandomForestClassifier}
  import org.apache.spark.ml.feature.{IndexToString, StringIndexer, VectorIndexer}
  import\ org. apache. spark. ml. evaluation. Binary Classification Evaluator
  import org.apache.spark.mllib.stat.{MultivariateStatisticalSummary, Statistics}
  import org.apache.spark.ml.evaluation.MulticlassClassificationEvaluator
  import org.apache.spark.ml.tuning.{CrossValidator, ParamGridBuilder}
  import\ org. a pache. spark. ml. classification. \{GBTClassification Model,\ GBTClassifier\}
  import\ org. apache. spark.ml. classification. Decision Tree Classification Model
  import org.apache.spark.ml.classification.DecisionTreeClassifier
```

```
// 1. Exploratory Data Analysis
// 1.1 Overview
// PassengerId is the unique id of the row and it doesn't have any effect on target
// Survived is the target variable we are trying to predict (0 or 1):
// 1 = Survived
// 0 = Not Survived
// Pclass (Passenger Class) is the socio-economic status of the passenger and it is a categorical ordinal feature which has 3
unique values (1, 2 or 3):
// 1 = Upper Class
// 2 = Middle Class
// 3 = Lower Class
// Name, Sex and Age are self-explanatory
// SibSp is the total number of the passengers' siblings and spouse
// Parch is the total number of the passengers' parents and children
// Ticket is the ticket number of the passenger
// Fare is the passenger fare
// Cabin is the cabin number of the passenger
// Embarked is port of embarkation and it is a categorical feature which has 3 unique values (C, Q or S):
// C = Cherbourg
// Q = Queenstown
// S = Southampton
// creating the schema for importing training and testing the datasets
val newTrainSchema = (new StructType)
.add("PassengerId", IntegerType)
.add("Survived", IntegerType)
.add("Pclass", IntegerType)
.add("Name", StringType)
.add("Sex", StringType)
.add("Age", FloatType)
.add("SibSp", IntegerType)
.add("Parch", IntegerType)
.add("Ticket", StringType)
.add("Fare", FloatType)
.add("Cabin", StringType)
.add("Embarked", StringType)
val newTestSchema = (new StructType)
.add("PassengerId", IntegerType)
.add("Pclass", IntegerType)
.add("Name", StringType)
.add("Sex", StringType)
.add("Age", FloatType)
.add("SibSp", IntegerType)
.add("Parch", IntegerType)
.add("Ticket", StringType)
.add("Fare", FloatType)
.add("Cabin", StringType)
.add("Embarked", StringType)
val trainSchema = StructType(newTrainSchema)
val testSchema = StructType(newTestSchema)
val csvFormat = "com.databricks.spark.csv"
val df_train = sqlContext.read.format(csvFormat).option("header","true").schema(trainSchema).load("/FileStore/tables/train-
4.csv")
val df_test = sqlContext.read.format(csvFormat).option("header","true").schema(testSchema).load("/FileStore/tables/test-5.csv")
//Creating table views for training and testing
df train.createOrReplaceTempView("df train")
df test.createOrReplaceTempView("df test")
```

newTrainSchema: org.apache.spark.sql.types.StructType = StructType(StructField(PassengerId,IntegerType,true),StructField(Survive d,IntegerType,true),StructField(Pass,IntegerType,true),StructField(Name,StringType,true),StructField(Sex,StringType,true),StructField(Age,FloatType,true),StructField(SibSp,IntegerType,true),StructField(Parch,IntegerType,true),StructField(Embarked,StringType,true),StructField(Embarked,StringType,true),StructField(Cabin,StringType,true),StructField(Embarked,StringType,true))

newTestSchema: org.apache.spark.sql.types.StructType = StructType(StructField(PassengerId,IntegerType,true),StructField(Pclass,I ntegerType,true),StructField(Age,FloatType,true),StructField(SibSp,IntegerType,true),StructField(Farch,IntegerType,true),StructField(Farch,IntegerType,true),StructField(Farch,FloatType,true),StructField(Cabin,StringType,true),StructField(Farch,FloatType,true))

trainSchema: org.apache.spark.sql.types.StructType = StructType(StructField(PassengerId,IntegerType,true),StructField(Survived,I ntegerType,true),StructField(Palsas,IntegerType,true),StructField(Name,StringType,true),StructField(Sex,StringType,true),StructField(ResperType,true),StructField(Ficket,StringType,true),StructField(Ficket,StringType,true),StructField(Ficket,StringType,true),StructField(Ficket,StringType,true),StructField(Ficket,StringType,true))

testSchema: org.apache.spark.sql.types.StructType = StructType(StructField(PassengerId,IntegerType,true),StructField(Pclass,IntegerType,true),StructField(Sib Sp,IntegerType,true),StructField(Sib Sp,IntegerType,true),StructField(

csvFormat: String = com.databricks.spark.csv

df_train: org.apache.spark.sql.DataFrame = [PassengerId: int, Survived: int ... 10 more fields]
df_test: org.apache.spark.sql.DataFrame = [PassengerId: int, Pclass: int ... 9 more fields]

```
//compute numcolumn summary statistics.
df_train.describe("Age", "SibSp", "Parch", "Fare").show()
```

summ	nary	Age	SibSp	Parch	Fare
+ cc	ount	714	891	891	891
_ m	nean 29.699	11764704046 0.5230	078563411896 0.38159	371492704824 32.20	9420804114722
sto	dev 14.5264	97332370992 1.1027	434322934315 0.8060	572211299488 49.69	9342916316158
1	min	0.42	0	0	0.0
1	max	80.0	8	6	512.3292

```
df_train.show()
```

PassengerId Sur	vived Pc				Ticket		Cabin Em	barked	
+	el	3 Braund, Mr. Owen male 22.0		+ 0	A/5 21171		+ null	- SI	
2	1	1 Cumings, Mrs. Joh female 38.0		øl	PC 17599			cl	
3	1	3 Heikkinen, Miss female 26.0		Ø STO	N/02. 3101282		0.00	sl	
41	1	1 Futrelle, Mrs. Ja female 35.0	1	0	113803	53.1	C123	sl	
5	0	3 Allen, Mr. Willia male 35.0	0	0	373450	8.05	null	sl	
6	0	3 Moran, Mr. James male null	0	0	330877	8.4583	null	QI	
7	0	1 McCarthy, Mr. Tim male 54.0	0	0	17463	51.8625	E46	sl	
8	0	3 Palsson, Master male 2.0	3	1	349909	21.075	null	sl	
9	1	3 Johnson, Mrs. Osc female 27.0	0	2	347742	11.1333	null	sl	
10	1	2 Nasser, Mrs. Nich female 14.0	1	0	237736	30.0708	null	cl	
11	1	3 Sandstrom, Miss female 4.0	1	1	PP 9549	16.7	G6	s	
12	1	1 Bonnell, Miss. El female 58.0	0	0	113783	26.55	C103	sl	
13	0	3 Saundercock, Mr male 20.0	0	0	A/5. 2151	8.05	null	sl	
14	0	3 Andersson, Mr. An male 39.0	1	5	347082	31.275	null	S	
15	0	3 Vestrom, Miss. Hu female 14.0	0	0	350406	7.8542	null	sl	-
16	1	2 Hewlett, Mrs. (Ma female 55.0	0	0	248706	16.0	null	s	•
17	0	3 Rice, Master. Eugene male 2.0	4	1	382652	29.125	null	QI	
18	1	2 Williams, Mr. Cha male null	0	0	244373	13.0	null	sl	//

//Summary stats for cateogrical columns sqlContext.sql("select Survived, count(*) from df_train group by Survived").show()

```
sqlContext.sql("select Pclass, Survived, count(*) from df_train group by Pclass, Survived").show()
```

 $sqlContext.sql("select Sex, Survived, count(*) from df_train group by Sex, Survived").show()$

```
//Calculating avg Age and Fare to fill null values for training
val AvgAge = df_train.select("Age")
    .agg(avg("Age"))
.collect() match {
    case Array(Row(avg: Double)) => avg
    case _ => 0
}
```

AvgAge: Double = 29.69911764704046

```
//Calculating the average fare for filling gaps in dataset train
val AvgFare = df_train.select("Fare")
  .agg(avg("Fare"))
  .collect() match {
    case Array(Row(avg: Double)) => avg
    case _ => 0
}
```

AvgFare: Double = 32.20420804114722

```
//Calculate avg Age and Fare to fill null values for test data
val AvgAge_test = df_test.select("Age")
  .agg(avg("Age"))
  .collect() match {
  case Array(Row(avg: Double)) => avg
  case _ => 0
}
```

AvgAge_test: Double = 30.272590361490668

```
//Calculate average fare for filling gaps in dataset test
val AvgFare_test = df_test.select("Fare")
    .agg(avg("Fare"))
    .collect() match {
    case Array(Row(avg: Double)) => avg
    case _ => 0
}
```

AvgFare_test: Double = 35.62718864996656

```
// for training
val embarked: (String => String) = {
  case "" => "S"
  case null => "S"
  case a => a
}
val embarkedUDF = udf(embarked)
```

embarked: String => String = \$Lambda\$8107/1598529161@41d4b905
embarkedUDF: org.apache.spark.sql.expressions.UserDefinedFunction = SparkUserDefinedFunction(\$Lambda\$8107/1598529161@41d4b905,String
Type,List(Some(class[value[0]: string])),Some(class[value[0]: string]),None,true,true)

```
//for test
val embarked_test: (String => String) = {
  case "" => "S"
  case null => "S"
  case a => a
}
val embarkedUDF_test = udf(embarked_test)
```

embarked_test: String => String = \$Lambda\$8119/1680966287@7f24d46
embarkedUDF_test: org.apache.spark.sql.expressions.UserDefinedFunction = SparkUserDefinedFunction(\$Lambda\$8119/1680966287@7f24d46,St
ringType,List(Some(class[value[0]: string])),Some(class[value[0]: string]),None,true,true)

```
//Filling null values with avg values for training dataset
val imputeddf = df_train.na.fill(Map("Fare" -> AvgFare, "Age" -> AvgAge))
val imputeddf2 = imputeddf.withColumn("Embarked", embarkedUDF(imputeddf.col("Embarked")))

//splitting training data into training and validation
val Array(trainingData, validationData) = imputeddf2.randomSplit(Array(0.7, 0.3))

//Filling null values with avg values for test dataset
val imputeddf_test = df_test.na.fill(Map("Fare" -> AvgFare_test, "Age" -> AvgAge_test))
val imputeddf2_test = imputeddf_test.withColumn("Embarked", embarkedUDF_test(imputeddf_test.col("Embarked")))
```

imputeddf: org.apache.spark.sql.DataFrame = [PassengerId: int, Survived: int ... 10 more fields]
imputeddf2: org.apache.spark.sql.DataFrame = [PassengerId: int, Survived: int ... 10 more fields]
trainingData: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [PassengerId: int, Survived: int ... 10 more fields]
validationData: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [PassengerId: int, Survived: int ... 10 more fields]
imputeddf_test: org.apache.spark.sql.DataFrame = [PassengerId: int, Pclass: int ... 9 more fields]
imputeddf2_test: org.apache.spark.sql.DataFrame = [PassengerId: int, Pclass: int ... 9 more fields]

```
// Feature Engineering - Create new attributes that may be derived from the existing attributes. This may include removing
 certain columns in the dataset.
  //Dropping Cabin feature as it has so many null values
  // val df1_train = trainingData.drop("Cabin")
  // val df1_test = imputeddf2_test.drop("Cabin")
 import org.apache.spark.ml.feature.Bucketizer
  import org.apache.spark.sql.functions._
  import org.apache.spark.sql.types._
  // Define the bucketizer splits based on the quantiles
  val splits = Array(Double.NegativeInfinity, 7.55, 7.854, 8.05, 10.5, 14.4542, 21.6792, 27.0, 39.6875, 77.9583, 512.329,
 Double.PositiveInfinity)
 // Create a Bucketizer transformer
  val bucketizer = new Bucketizer()
    .setInputCol("Fare") // Specify the input column
   .setOutputCol("FareBucket") // Specify the output column
 // Apply the bucketizer transformation to the {\tt DataFrame}
  val df_binned = bucketizer.transform(df_train)
 // Now df_binned contains the 'FareBucket' column, which represents the binned 'Fare' values
{\tt import\ org.apache.spark.ml.feature.Bucketizer}
import org.apache.spark.sql.functions._
import org.apache.spark.sql.types._
splits: Array[Double] = Array(-Infinity, 7.55, 7.854, 8.05, 10.5, 14.4542, 21.6792, 27.0, 39.6875, 77.9583, 512.329, Infinity)
```

df_binned.show()

+		+	+	+	+		+	+	+	+	
Pass	engerId Sur	vived Pc	lass Name Sex Age S	bSp Pa	rch	Ticket	Fare	Cabin Emb	arked Far	eBucket	
+		+			anne a gant an					+	۰
I	1	0	3 Braund, Mr. Owen male 22.0	1	0	A/5 21171		null	s	0.0	-1
	2	1	1 Cumings, Mrs. Joh female 38.0	1	0	PC 17599	71.2833	C85	c	8.0	-1
	3	1	3 Heikkinen, Miss female 26.0	0	0 STC	N/02. 3101282	7.925	null	S	2.0	-1
	4	1	1 Futrelle, Mrs. Ja female 35.0	1	0	113803	53.1	C123	S	8.0	-
	5	0	3 Allen, Mr. Willia male 35.0	0	0	373450	8.05	null	s	3.0	-1
	6	0	3 Moran, Mr. James male null	0	0	330877	8.4583	null	Ql	3.0	-1
	7	0	1 McCarthy, Mr. Tim male 54.0	0	0	17463	51.8625	E46	s	8.0	- 1
	8	0	3 Palsson, Master male 2.0	3	1	349909	21.075	null	s	5.0	- 1
	9	1	3 Johnson, Mrs. Osc female 27.0	0	2	347742	11.1333	null	s	4.0	- 1
	10	1	2 Nasser, Mrs. Nich female 14.0	1	0	237736	30.0708	null	cl	7.0	
	11	1	3 Sandstrom, Miss female 4.0	1	1	PP 9549	16.7	G6	s	5.0	- 1
	12	1	1 Bonnell, Miss. El female 58.0	0	0	113783	26.55	C103	s	6.0	
	13	0	3 Saundercock, Mr male 20.0	0	0	A/5. 2151	8.05	null	s	3.0	
	14	0	3 Andersson, Mr. An male 39.0	1	5	347082	31.275	null	s	7.0	
	15	0	3 Vestrom, Miss. Hu female 14.0	0	0	350406	7.8542	null	s	2.0	
	16	1	2 Hewlett, Mrs. (Ma female 55.0	0	0	248706	16.0	null	sl	5.0	
	17	0	3 Rice, Master. Eugene male 2.0	4	1	382652	29.125	null	QI	7.0	
	18	1	2 Williams, Mr. Cha male null	0	0	244373	13.0	null	sl	4.0	

bucketizer: org.apache.spark.ml.feature.Bucketizer = Bucketizer: uid=bucketizer_4leda3470974 df_binned: org.apache.spark.sql.DataFrame = [PassengerId: int, Survived: int ... 11 more fields]

```
Splits: Array[Double] = Array(-Infinity, 19.0, 22.0, 25.0, 28.0, 31.0, 35.0, 39.0, 45.0, 55.0, 65.0, Infinity) bucketizer: org.apache.spark.ml.feature.Bucketizer = Bucketizer: uid=bucketizer_55340fd47ead df_binned: org.apache.spark.sql.DataFrame = [PassengerId: int, Survived: int ... 11 more fields]
```

df_binned.show()

PassengerId Sur	vived Pc	lass Name Sex Age S			Ticket		Cabin Emb		eBucket	
1	0	3 Braund, Mr. Owen male 22.0	1	0	A/5 21171		null	s	2.0	
2	1	1 Cumings, Mrs. Joh female 38.0	1	0	PC 17599	71.2833	C85	cl	6.0	
3	1	3 Heikkinen, Miss female 26.0	0	0 ST	ON/O2. 3101282	7.925	null	s	3.0	
4	1	1 Futrelle, Mrs. Ja female 35.0	1	0	113803	53.1	C123	sl	6.0	
5	0	3 Allen, Mr. Willia male 35.0	0	0	373450	8.05	null	sl	6.0	
6	0	3 Moran, Mr. James male null	0	0	330877	8.4583	null	QI	null	
7	0	1 McCarthy, Mr. Tim male 54.0	0	0	17463	51.8625	E46	sl	8.0	-
8	0	3 Palsson, Master male 2.0	3	1	349909	21.075	null	S	0.0	-
9	1	3 Johnson, Mrs. Osc female 27.0	0	2	347742	11.1333	null	s	3.0	-
10	1	2 Nasser, Mrs. Nich female 14.0	1	0	237736	30.0708	null	c	0.0	
11	1	3 Sandstrom, Miss female 4.0	1	1	PP 9549	16.7	G6	s	0.0	
12	1	1 Bonnell, Miss. El female 58.0	0	0	113783	26.55	C103	S	9.0	
13	0	3 Saundercock, Mr male 20.0	0	0	A/5. 2151	8.05	null	sl	1.0	
14	0	3 Andersson, Mr. An male 39.0	1	5	347082	31.275	null	s	7.0	
15	0	3 Vestrom, Miss. Hu female 14.0	0	0	350406	7.8542	null	sl	0.0	
16	1	2 Hewlett, Mrs. (Ma female 55.0	0	0	248706	16.0	null	sl	9.0	
17	0	3 Rice, Master. Eugene male 2.0	4	1	382652	29.125	null	QI	0.0	
18	1	2 Williams, Mr. Cha male null	el	øl	244373	13.0	null	sl	null	

```
import org.apache.spark.sql.functions._
import org.apache.spark.sql.DataFrame

// Adding a new column 'Family_Size' to df_all
val df_with_family_size: DataFrame = df_binned.withColumn("Family_Size", col("SibSp") + col("Parch") + lit(1))

// Now df_with_family_size contains the 'Family_Size' column
```

import org.apache.spark.sql.functions._
import org.apache.spark.sql.DataFrame
df_with_family_size: org.apache.spark.sql.DataFrame = [PassengerId: int, Survived: int ... 12 more fields]

df with family size.show() |PassengerId|Survived|Pclass| Name | Sex | Age | SibSp | Parch | Ticket| Fare|Cabin|Embarked|AgeBucket|Fam ily_Size| ----+ 3|Braund, Mr. Owen ...| male|22.0| 1| 0| A/5 21171 7.25 | null| 1 0 S 2.0 2| 2 1 1|Cumings, Mrs. Joh...|female|38.0| 1| 0| PC 17599 | 71.2833 | C85 | cl 6.0 2 3| 11 3|Heikkinen, Miss. ...|female|26.0| 0| 0|STON/02. 3101282| 7.925| null| sl 3.0 1 41 1 1|Futrelle, Mrs. Ja...|female|35.0| 1| 113803| 53.1| C123| SI 6.0 21 3|Allen, Mr. Willia...| male|35.0| 373450| 8.05| null| 5 0 0 01 SI 6.0 1 6 0 3| Moran, Mr. James| male|null| 330877| 8.4583| null| Ql null| 11 71 0 1|McCarthy, Mr. Tim...| male|54.0| 0| 17463 | 51.8625 | E46 | SI 0 8.0 1 8 0 3|Palsson, Master. ...| male| 2.0| 3| 1| 349909| 21.075| null| import org.apache.spark.sql.expressions.Window import org.apache.spark.sql.functions._ import org.apache.spark.sql.DataFrame // Define a Window specification to partition by 'Ticket' column val windowSpec = Window.partitionBy("Ticket") // Add a new column 'Ticket_Frequency' to df_all $val \ df_with_ticket_frequency: \ DataFrame = df_with_family_size.withColumn("Ticket_Frequency", \ count("*").over(windowSpec))$ // Now df with ticket frequency contains the 'Ticket Frequency' column import org.apache.spark.sql.expressions.Window import org.apache.spark.sql.functions._ import org.apache.spark.sql.DataFrame windowSpec: org.apache.spark.sql.expressions.WindowSpec = org.apache.spark.sql.expressions.WindowSpec@3e8679a5 df_with_ticket_frequency: org.apache.spark.sql.DataFrame = [PassengerId: int, Survived: int ... 13 more fields] df_with_ticket_frequency.show() |PassengerId|Survived|Pclass| Name | Sex | Age | SibSp | Parch | Ticket | Fare | Cabin | Embarked | AgeBucket | Family_Size | T icket_Frequency| ----+ 1 1|Cherry, Miss. Gladys|female|30.0| 0| 0|110152| 86.5| B77| S 1 258 4.0

0|110152| 86.5| B79|

0|110152| 86.5| B77|

1|110413| 79.65| E67|

S

S

0.0

5.0

8.0

1

1|

3|

505

760

263

3

1

1

0

1|Maioni, Miss. Rob...|female|16.0| 0|

1|Rothes, the Count...|female|33.0| 0|

1| Taussig, Mr. Emil| male|52.0| 1|

```
559
                   1
                          1|Taussig, Mrs. Emi...|female|39.0| 1|
                                                                     1|110413| 79.65| E67|
                                                                                                    S
                                                                                                            7.0
                                                                                                                          3|
3
                          1 | Taussig, Miss. Ruth | female | 18.0 |
                                                                      2|110413| 79.65| E68|
                                                                                                                          3|
         586
                   1
                                                                                                    SI
                                                                                                            0.0
         111
                   0
                          1|Porter, Mr. Walte...| male|47.0| 0|
                                                                      0|110465| 52.0| C110|
                                                                                                            8.0|
                                                                                                                          1|
2|
  import org.apache.spark.sql.{DataFrame.SparkSession}
  import org.apache.spark.sql.functions.
  // Initialize a SparkSession
  val spark = SparkSession.builder()
    .appName("DataFrame Transformation")
    .getOrCreate()
  // Add a new column 'Title' by splitting the 'Name' column
  val df_with_title: DataFrame = df_with_ticket_frequency.withColumn("Title", split(col("Name"), ", ")(1))
    .withColumn("Title", split(col("Title"), "\\.")(0))
  // Add a new column 'Is_Married' and set it to 1 for 'Title' equal to 'Mrs'
  val df_with_is_married: DataFrame = df_with_title.withColumn("Is_Married", when(col("Title") === "Mrs", 1).otherwise(0))
  // Now df_with_is_married contains the 'Title' and 'Is_Married' columns
import org.apache.spark.sql.{DataFrame, SparkSession}
import org.apache.spark.sql.functions._
spark: org.apache.spark.sql.SparkSession = org.apache.spark.sql.SparkSession@3d025628
df_with_title: org.apache.spark.sql.DataFrame = [PassengerId: int, Survived: int ... 14 more fields]
df_with_is_married: org.apache.spark.sql.DataFrame = [PassengerId: int, Survived: int ... 15 more fields]
```

df_with_is_married.show() Name| Sex| Age|SibSp|Parch|Ticket| Fare|Cabin|Embarked|AgeBucket|Family_Size|T |PassengerId|Survived|Pclass| Title|Is Married| icket Frequency 258 1 1|Cherry, Miss. Gladys|female|30.0| 0| 0|110152| 86.5| B77| SI 4.0 1 3 Miss 1 1|Maioni, Miss. Rob...|female|16.0| 0| 0|110152| 86.5| B79| 505 S 0.0 1| 3| Miss| 1 1|Rothes, the Count...|female|33.0| 0| 7601 0|110152| 86.5| B77| S 5.0 1 0 3|the Countess| 263 0 1| Taussig, Mr. Emil| male|52.0| 1| 1|110413| 79.65| E67| s 8.0 3 3| 1 1|Taussig, Mrs. Emi...|female|39.0| 1| 1|110413| 79.65| E67| 559 S 7.0 3 31 1 Mrs 1 1| Taussig, Miss. Ruth|female|18.0| 586 2|110413| 79.65| E68| s 0.0 3 Miss 111 0 1|Porter, Mr. Walte...| male|47.0| 0|110465| 52.0| C110| S 8.0 11 01 2 Mr 0| 1|Clifford, Mr. Geo...| male|null| 0| 0|110465| 52.0| A14| 476 null| 1| S

```
import org.apache.spark.sql.functions._
  import org.apache.spark.sql.{DataFrame, SparkSession}
  import java.util.regex.Pattern
   import java.util.regex.Matcher
  // Assuming you have a DataFrame df_all
  // Initialize a SparkSession
  val spark = SparkSession.builder()
    .appName("DataFrame Transformation")
    .getOrCreate()
  // Define a user-defined function (UDF) to extract surnames
  val extractSurname = udf((name: String) => {
    val nameNoBracket = if (name.contains("(")) name.split("\\(")(0) else name
    val family = nameNoBracket.split(",")(0).replaceAll("[^a-zA-Z]", "").trim
    family
  })
  \ensuremath{//} Add a new column 'Family' by applying the UDF to the 'Name' column
  val df_with_family: DataFrame = df_with_is_married.withColumn("Family", extractSurname(col("Name")))
  // Now df_train and df_test are the training and testing DataFrames, and df_with_family contains the 'Family' column
import org.apache.spark.sql.functions._
import org.apache.spark.sql.{DataFrame, SparkSession}
import java.util.regex.Pattern
import java.util.regex.Matcher
spark: org.apache.spark.sql.SparkSession = org.apache.spark.sql.SparkSession@3d025628
extractSurname: org.apache.spark.sql.expressions.UserDefinedFunction = SparkUserDefinedFunction($Lambda$8827/125107648@b431115,Strin
gType,List(Some(class[value[0]: string])),Some(class[value[0]: string]),None,true,true)
df_with_family: org.apache.spark.sql.DataFrame = [PassengerId: int, Survived: int ... 16 more fields]
df_train: org.apache.spark.sql.DataFrame = [PassengerId: int, Survived: int ... 16 more fields]
df_test: org.apache.spark.sql.DataFrame = [PassengerId: int, Survived: int ... 16 more fields]
```

df_with_family.show()

Passe	ngerId Survi	ived P	class	Name S	ex Age Si	bSp Pa	rch Ticket	Fare	Cabin Emb	arked Ag	eBucket Fami	ly_Size T
icket_	Frequency		Title Is_Married		Family							
							+					
I	258	1	1 Cherry, Miss	s. Gladys fema	le 30.0	0	0 110152	86.5	B77	s	4.0	1
3	Miss		0	herry								
1	505	1	1 Maioni, Miss	s. Rob fema	le 16.0	0	0 110152	86.5	B79	s	0.0	1
3	Miss		0 1	Maioni								
1	760	1	1 Rothes, the	Count fema	le 33.0	0	0 110152	86.5	B77	s	5.0	1
3 the	Countess		0	Rothes								
1	263	0	1 Taussig,	Mr. Emil ma	le 52.0	1	1 110413	79.65	E67	s	8.0	3
3	Mr		0 Ta	aussig								
I	559	1	1 Taussig, Mrs	s. Emi fema	le 39.0	1	1 110413	79.65	E67	s	7.0	3
3	Mrs		1 Ta	aussig								
[586	1	1 Taussig, M	iss. Ruth fema	le 18.0	0	2 110413	79.65	E68	s	0.0	3
3	Miss		0 Ta	aussig								
1	111	0	1 Porter, Mr.	Walte ma	le 47.0	0	0 110465	52.0	C110	S	8.0	1
2	Mr		0	Porter								
1	476	0	1 Clifford, Mr	. Geo ma	le null	0	0 110465	52.0	A14	S	null	1

```
//Indexing categorical features
    val catFeatColNames = Seq("Pclass", "Sex", "Embarked")
    val stringIndexers = catFeatColNames.map { colName =>
       new StringIndexer()
            .setInputCol(colName)
             .setOutputCol(colName + "Indexed")
            .fit(trainingData)
    //Indexing target feature
    val labelIndexer = new StringIndexer()
    .setInputCol("Survived")
    .setOutputCol("SurvivedIndexed")
    .fit(trainingData)
    //Assembling features into one vector
    val numFeatColNames = Seq("Age", "SibSp", "Parch", "Fare")
    val idxdCatFeatColName = catFeatColNames.map(_ + "Indexed")
    val allIdxdFeatColNames = numFeatColNames ++ idxdCatFeatColName
    val assembler = new VectorAssembler()
       .setInputCols(Array(allIdxdFeatColNames: _*))
        .setOutputCol("Features")
catFeatColNames: Seq[String] = List(Pclass, Sex, Embarked)
stringIndexers: Seq[org.apache.spark.ml.feature.StringIndexerModel] = List(StringIndexerModel: uid=strIdx_52c20ddf3e58, handleInvali
d=error, StringIndexerModel: uid=strIdx_510c7847544f, handleInvalid=error, StringIndexerModel: uid=strIdx_16154abc0509, handleInvali
label Indexer: org. apache. spark. \verb|m|l.feature.StringIndexerModel| = StringIndexerModel|: uid=strIdx\_d082391f2c11, handle Invalid=error apache. StringIndexerModel| = String
numFeatColNames: Seq[String] = List(Age, SibSp, Parch, Fare)
idxdCatFeatColName: Seq[String] = List(PclassIndexed, SexIndexed, EmbarkedIndexed)
all IdxdFeatColNames: Seq[String] = List(Age, SibSp, Parch, Fare, PclassIndexed, SexIndexed, EmbarkedIndexed)\\
assembler: org.apache.spark.ml.feature.VectorAssembler = VectorAssembler: uid=vecAssembler_fe05be893203, handleInvalid=error, numInp
utCols=7
    //Training the model from the pool and determining the best one among them.
    //Randomforest classifier
    val randomforest = new RandomForestClassifier()
       .setLabelCol("SurvivedIndexed")
        .setFeaturesCol("Features")
    //Retrieving original labels
    val labelConverter = new IndexToString()
       .setInputCol("prediction")
        .setOutputCol("predictedLabel")
        .setLabels(labelIndexer.labelsArray.flatten)
    //Creating pipeline
    val pipeline = new Pipeline().setStages(
       (stringIndexers :+ labelIndexer :+ assembler :+ randomforest :+ labelConverter).toArray)
```

randomforest: org.apache.spark.ml.classification.RandomForestClassifier = rfc_d39929b5363e labelConverter: org.apache.spark.ml.feature.IndexToString = idxToStr_deec7f48ee96 pipeline: org.apache.spark.ml.Pipeline = pipeline_f7c61e980795

```
//Selecting best model
  val paramGrid = new ParamGridBuilder()
    .addGrid(randomforest.maxBins, Array(25, 28, 31))
    .addGrid(randomforest.maxDepth, Array(4, 6, 8))
    .addGrid(randomforest.impurity, Array("entropy", "gini"))
    .build()
  val evaluator = new BinaryClassificationEvaluator()
    .setLabelCol("SurvivedIndexed")
    .setMetricName("areaUnderPR")
  //Cross validator with 10 fold
  val cv = new CrossValidator()
    .setEstimator(pipeline)
    .setEvaluator(evaluator)
    . {\tt setEstimatorParamMaps} ({\tt paramGrid})
    .setNumFolds(10)
paramGrid: Array[org.apache.spark.ml.param.ParamMap] =
        rfc_d39929b5363e-impurity: entropy,
        rfc_d39929b5363e-maxBins: 25,
        rfc_d39929b5363e-maxDepth: 4
        rfc_d39929b5363e-impurity: gini,
        rfc_d39929b5363e-maxBins: 25,
        rfc_d39929b5363e-maxDepth: 4
        rfc_d39929b5363e-impurity: entropy,
        rfc d39929b5363e-maxBins: 25,
        rfc_d39929b5363e-maxDepth: 6
}, {
        rfc_d39929b5363e-impurity: gini,
        rfc d39929b5363e-maxBins: 25,
        rfc_d39929b5363e-maxDepth: 6
        rfc_d39929b5363e-impurity: entropy,
        rfc d39929b5363e-maxBins: 25.
        rfc_d39929b5363e-maxDepth: 8
  //Fitting model using cross validation
  val crossValidatorModel = cv.fit(trainingData)
  //predictions on validation data
  val predictions = crossValidatorModel.transform(validationData)
  //Accuracy
  val accuracy = evaluator.evaluate(predictions)
  println("Test Error DT= " + (1.0 - accuracy))
Test Error DT= 0.19512230928774577
crossValidatorModel: org.apache.spark.ml.tuning.CrossValidatorModel = CrossValidatorModel: uid=cv_ce476272474c, bestModel=pipeline_f
7c61e980795, numFolds=10
predictions: org.apache.spark.sql.DataFrame = [PassengerId: int, Survived: int ... 19 more fields]
accuracy: Double = 0.8048776907122542
Test Error = 0.20172748564799914
gbt: org.apache.spark.ml.classification.GBTClassifier = gbtc_145212851005
pipeline: org.apache.spark.ml.Pipeline = pipeline_200b998877e7
model: org.apache.spark.ml.PipelineModel = pipeline_200b998877e7 predictions: org.apache.spark.sql.DataFrame = [PassengerId: int, Survived: int ... 19 more fields]
evaluator: org.apache.spark.ml.evaluation.BinaryClassificationEvaluator = BinaryClassificationEvaluator: uid=binEval_913ee5288d5a, m
etricName=areaUnderPR, numBins=1000
accuracy: Double = 0.7982725143520009
```